

# Q-learning Algorithm

Justin Sirignano

University of Illinois at Urbana-Champaign

October 17, 2019

# Deep Reinforcement Learning

- ▶ Algorithms for training deep learning models to estimate an optimal control.
- ▶ Successes: Atari games, Go, Starcraft.
- ▶ Potential applications: robotics, driverless cars, UAVs.
- ▶ Relatively few mathematical guarantees.

# Markov Decision Problem

- ▶ We consider a Markov decision problem defined on finite state space  $\mathcal{X} \subset \mathbb{R}^{d_x}$ .
- ▶ In each state, an action  $a \in \mathcal{A} \subset \mathbb{R}^{d_a}$  can be taken.
- ▶ The probability transition function is  $p(z|x, a) = \mathbb{P}[x_{j+1} = z | x_j = x, a_j = a]$ .
- ▶ The reward function is  $r(x, a)$ .
- ▶ The objective is to determine a policy  $a^* : \mathcal{X} \rightarrow \mathcal{A}$  that maximizes

$$W(x) = \mathbb{E} \left[ \sum_{j=0}^{\infty} \gamma^j r(x_j, a_j) | x_0 = x \right], \quad (1)$$

where  $a_j = a^*(x_j)$ .

# Bellman Equation

$$\begin{aligned} 0 &= r(x, a) + \gamma \sum_{z \in \mathcal{X}} \max_{a' \in \mathcal{A}} V(z, a') p(z|x, a) - V(x, a), \\ a^*(x) &= \arg \max_{a \in \mathcal{A}} V(x, a). \end{aligned} \tag{2}$$

- ▶ In principle, the Bellman equation can be solved to find the optimal control.
- ▶ However, the transition probability function  $p(z|x, a)$  (i.e., the state dynamics) may not be known.
- ▶ Even if they are known, the state space may be too high-dimensional for standard numerical methods to solve (2) due to the curse-of-dimensionality.

# Reinforcement Learning

- ▶ Reinforcement learning approximates the solution to the Bellman equation with a function approximator  $Q(x, a; \theta)$  such as a (deep) neural network.
- ▶ The parameters of the neural network, denoted by  $\theta$ , are estimated using the Q-learning algorithm.
- ▶ The neural network in Q-learning is referred to as a “Q-network”.

# Q-learning

The Q-learning algorithm attempts to minimize the objective function

$$L(\theta) = \sum_{x,a \in \mathcal{X} \times \mathcal{A}} \left[ (Y(x,a) - Q(x,a;\theta))^2 \right] \pi(x,a), \quad (3)$$

where  $\pi$  is a probability measure which is positive for every  $(x,a)$  and the "target"  $Y$  is

$$Y(x,a) = r(x,a) + \gamma \sum_{x' \in \mathcal{X}} \max_{a' \in \mathcal{A}} Q(x',a';\theta) p(x'|x,a). \quad (4)$$

- ▶ If  $L(\theta) = 0$ , then  $Q(x,a;\theta)$  is a solution to the Bellman equation.
- ▶ In practice, the hope is that the Q-learning algorithm will learn a model  $Q$  such that  $L(\theta)$  is small and therefore  $Q(x,a;\theta)$  is a good approximation for the Bellman solution  $V(x,a)$ .

$$L(\theta) = \sum_{x,a \in \mathcal{X} \times \mathcal{A}} \left[ (Y(x,a) - Q(x,a;\theta))^2 \right] \pi(x,a),$$

$$Y(x,a) = r(x,a) + \gamma \sum_{x' \in \mathcal{X}} \max_{a' \in \mathcal{A}} Q(x',a';\theta) p(x'|x,a). \quad (5)$$

- ▶ Take the “gradient” of  $L(\theta)$  but treat  $Y$  as a constant.

$$\theta_{k+1} = \theta_k - \alpha \sum_{x,a} \left[ (Y(x,a) - Q(x,a;\theta)) \right] \nabla_{\theta} Q(x,a;\theta) \pi(x,a). \quad (6)$$

- ▶ To make the update (6) computationally efficient, use the stochastic approximation:

$$\begin{aligned} \theta_{k+1} &= \theta_k + \alpha G_k, \\ G_k &= \left( r(x_k, a_k) + \gamma \max_{a' \in \mathcal{A}} Q(x_{k+1}, a'; \theta_k) - Q(x_k, a_k; \theta_k) \right) \\ &\quad \times \nabla_{\theta} Q(x_k, a_k; \theta_k), \end{aligned} \quad (7)$$

where, for example,  $(x_k, a_k)$  is an ergodic Markov chain with a unique stationary distribution  $\pi(x, a)$ .

## Deep Q-learning

A “Q-network”  $Q(x, a; \theta)$  is trained using the Q-learning algorithm

$$\begin{aligned}\theta_{k+1} &= \theta_k + \alpha G_k, \\ G_k &= \left( y_k - Q(x_k, a_k; \theta_k) \right) \nabla_{\theta} Q(x_k, a_k; \theta_k),\end{aligned}\tag{8}$$

where

$$y_k = r(x_k, a_k) + \gamma \max_{a' \in \mathcal{A}} Q(x_{k+1}, a'; \theta_k)\tag{9}$$

is considered to be the “target” for iteration  $k$ .

- ▶ Unlike in the stochastic gradient descent algorithm, the Q-learning update directions  $G_k$  are not unbiased estimates of a descent direction for the objective function  $L(\theta)$ .
- ▶ The Q-learning algorithm takes the derivative of  $L(\theta)$  *while treating the target  $Y$  as a constant*. Therefore,

$$\mathbb{E}[G_k | \theta_k, x_k, a_k] \neq \nabla_{\theta} \left[ (Y(x_k, a_k) - Q(x_k, a_k; \theta_k))^2 \right].\tag{10}$$



# Q-learning: mathematical challenges

- ▶ Unlike in the stochastic gradient descent algorithm, the Q-learning update directions  $G_k$  are not unbiased estimates of a descent direction for the objective function  $L(\theta)$ .
- ▶ It is not necessarily clear that an update step will lead to the objective function decreasing.
- ▶ Many of the mathematical challenges of analyzing Q-learning arise from this fact.

A  $Q$ -network is trained using the Q-learning algorithm

$$\theta_{k+1} = \theta_k + \left( y_k - Q(x_k, a_k; \theta_k) \right) \nabla_{\theta} Q(x_k, a_k; \theta_k), \quad (11)$$

where

$$y_k = r(x_k, a_k) + \gamma \max_{a' \in \mathcal{A}} Q(x_{k+1}, a'; \theta_k) \quad (12)$$

is considered to be the target for iteration  $k$ .

The Q-learning algorithm estimates an approximation to the Bellman equation using stochastic gradient descent (but where the gradient is not taken with respect to the target  $y_k$ ).