

PyTorch Tutorial

Rajbir Kataria

09/17/2018

Overview

- Tensors
- Optimizers
- nn module
- Examples! Examples! Examples!
 - https://github.com/yunjey/pytorch-tutorial/blob/master/tutorials/01-basics/pytorch_basics/main.py

PyTorch Tensors

- Similar to numpy's ndarrays
 - Can also be used on a GPU
 - Lost of operations supported like '+' and '-'

```
# import pytorch
import torch

# define a tensor
torch.FloatTensor([2])
```

```
2
[torch.FloatTensor of size 1]
```

Optim module

- Module that implements various optimization algorithms used for building neural networks
 - Most common methods are already supported

```
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

nn module

- Easy to define computational graphs and take gradients
 - nn package defines a set of modules that can be used to make a neural network

```
import torch

# define model
model = torch.nn.Sequential(
    torch.nn.Linear(input_num_units, hidden_num_units),
    torch.nn.ReLU(),
    torch.nn.Linear(hidden_num_units, output_num_units),
)
loss_fn = torch.nn.CrossEntropyLoss()
```

Loading data from numpy

```
# Create a numpy array.
```

```
x = np.array([[1, 2], [3, 4]])
```

```
# Convert the numpy array to a torch tensor.
```

```
y = torch.from_numpy(x)
```

```
# Convert the torch tensor to a numpy array.
```

```
z = y.numpy()
```

Autograd example

```
# Create tensors of shape (10, 3) and (10, 2).
x = torch.randn(10, 3)
y = torch.randn(10, 2)

# Build a fully connected layer.
linear = nn.Linear(3, 2)
print('w: ', linear.weight)
print('b: ', linear.bias)

# Build loss function and optimizer.
criterion = nn.MSELoss()
optimizer = torch.optim.SGD(linear.parameters(), lr=0.01)

# Forward pass.
pred = linear(x)

# Compute loss.
loss = criterion(pred, y)
print('loss: ', loss.item())

# Backward pass.
loss.backward()

# Print out the gradients.
print('dL/dw: ', linear.weight.grad)
print('dL/db: ', linear.bias.grad)

# 1-step gradient descent.
optimizer.step()
```

Input pipeline

```
# Download and construct CIFAR-10 dataset.
train_dataset = torchvision.datasets.CIFAR10(root='.././data/',
                                              train=True,
                                              transform=transforms.ToTensor(),
                                              download=True)

# Fetch one data pair (read data from disk).
image, label = train_dataset[0]

# Data loader (this provides queues and threads in a very simple way).
train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                             batch_size=64,
                                             shuffle=True)

# When iteration starts, queue and thread start to load data from files.
data_iter = iter(train_loader)

# Mini-batch images and labels.
images, labels = data_iter.next()

# Actual usage of the data loader is as below.
for images, labels in train_loader:
    # Training code should be written here.
    pass
```


GPU/CPU interaction

```
# Specify device to use
device = torch.device("cuda" if use_cuda else "cpu")

# Use the specific device
model = Net().to(device)

# Put data on the device and pass it to the model
for batch_idx, (data, target) in enumerate(train_loader):
    data, target = data.to(device), target.to(device)
    optimizer.zero_grad()
    output = model(data)

    # Convert output to CPU and print
    print (output.data.cpu().numpy())
```

Pretrained model

```
# Download and load the pretrained ResNet-18.
resnet = torchvision.models.resnet18(pretrained=True)

# If you want to finetune only the top layer of the model, set as below.
for param in resnet.parameters():
    param.requires_grad = False

# Replace the top layer for finetuning.
resnet.fc = nn.Linear(resnet.fc.in_features, 100) # 100 is an example.

# Forward pass.
images = torch.randn(64, 3, 224, 224)
outputs = resnet(images)
print (outputs.size())      # (64, 100)
```

Save and load the model

```
# Save and load the entire model.
```

```
torch.save(resnet, 'model.ckpt')
```

```
model = torch.load('model.ckpt')
```

```
# Save and load only the model parameters (recommended).
```

```
torch.save(resnet.state_dict(), 'params.ckpt')
```

```
resnet.load_state_dict(torch.load('params.ckpt'))
```

References

- All examples from:
 - https://github.com/yunjey/pytorch-tutorial/blob/master/tutorials/01-basics/pytorch_basics/main.py
- Other resources:
 - <https://www.analyticsvidhya.com/blog/2018/02/pytorch-tutorial/>